



*All 64 slices share the same backbone weights | Only the 256-dim MIL embedding reaches the classifier*